

# Chapter 2

## Application Layer

These slides have been adapted from the Pearson slides. S. Thomson 09/2017.

### A note on the use of these Powerpoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

 All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



## Computer Networking: A Top Down Approach

7<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson/Addison Wesley  
April 2016

# Chapter 2: outline



The image part with relationship ID rld3 was not found in the file.

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP


# Content

- What transport services does the network provide to applications?
- How does the application developer make use of these services?
- How does the application developer decide what service to use?

# RECAP: What is the Internet Architecture?

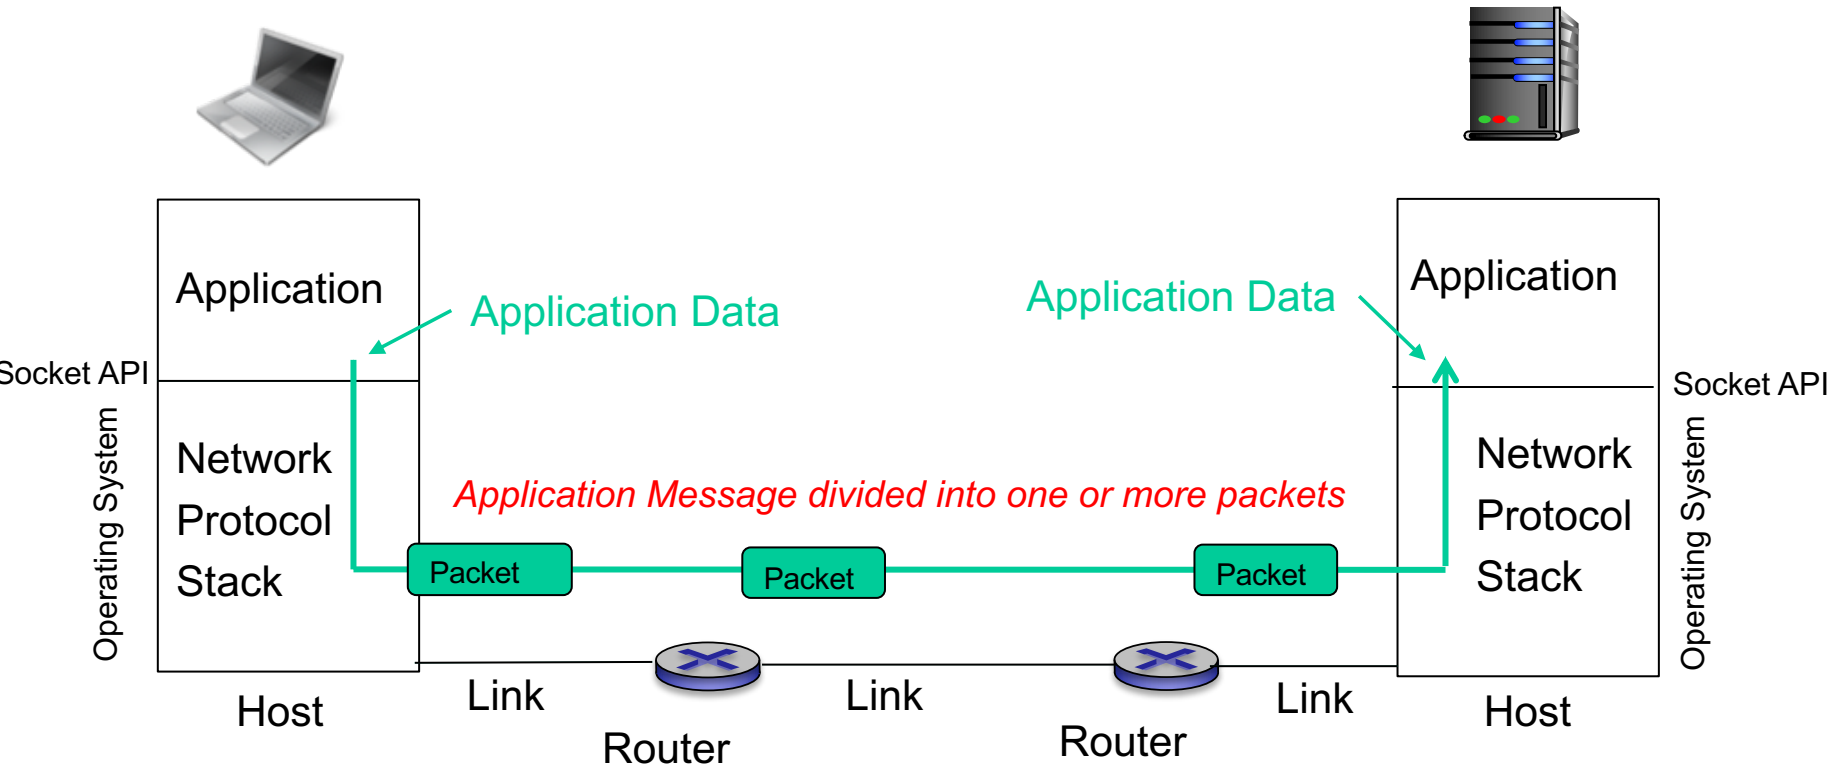
## Packet Switching

- Components:
  - Hosts: Transmit and receive application **packets** of length **L**
  - Routers: Store and **forward** packets from one link to another
  - Links: **Propagate** packets within the link
- Note:
  - A router may also act as a host e.g. management



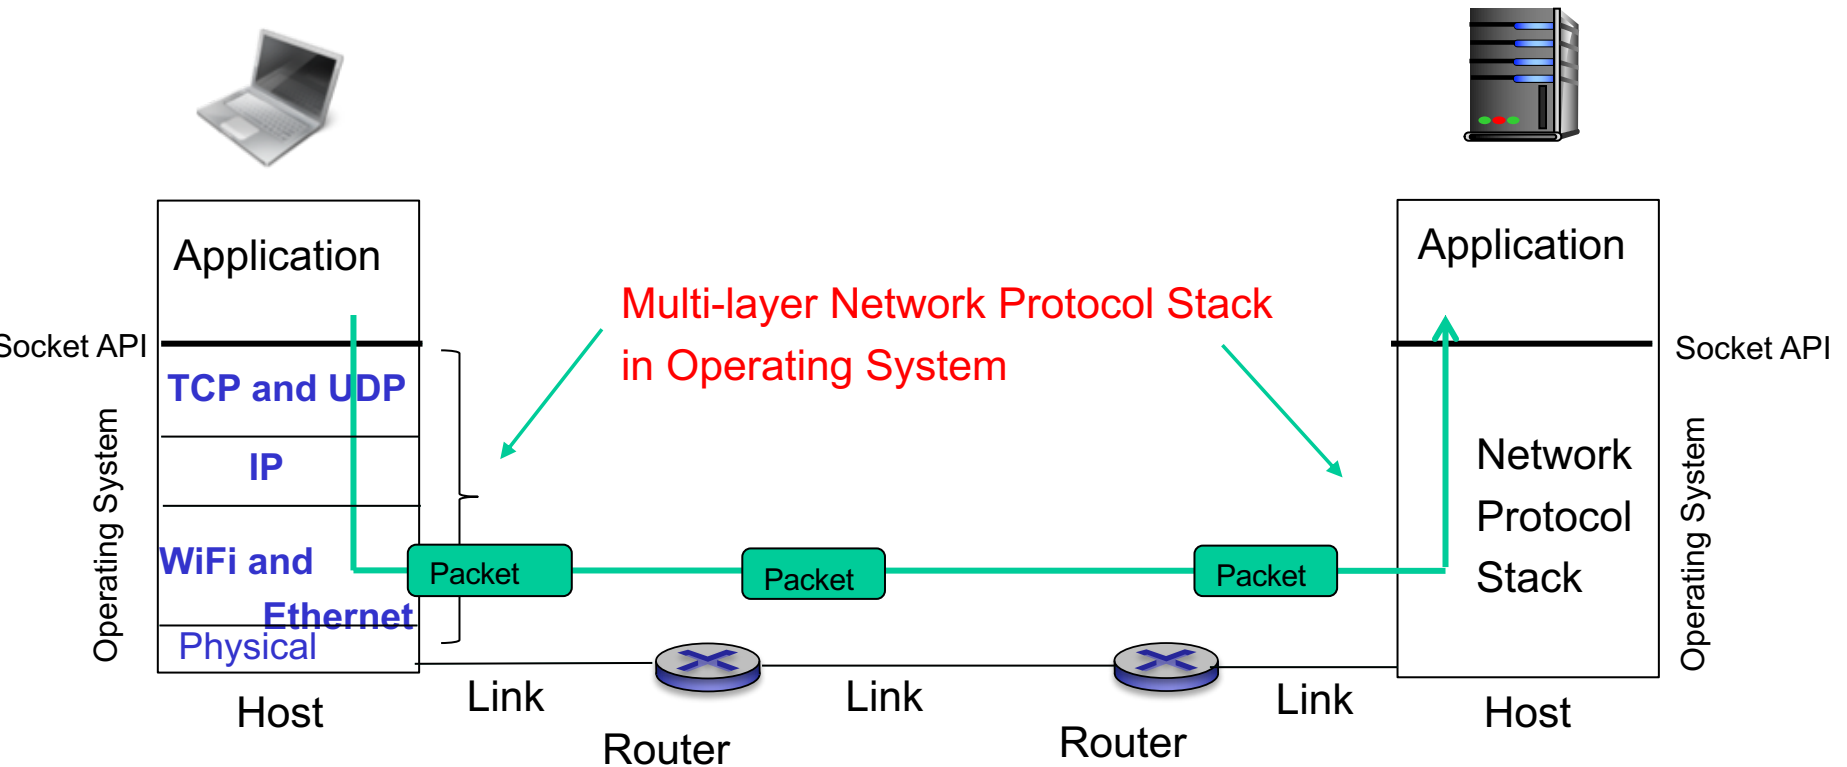
# RECAP: What is the Internet Architecture?

## Packet Switching



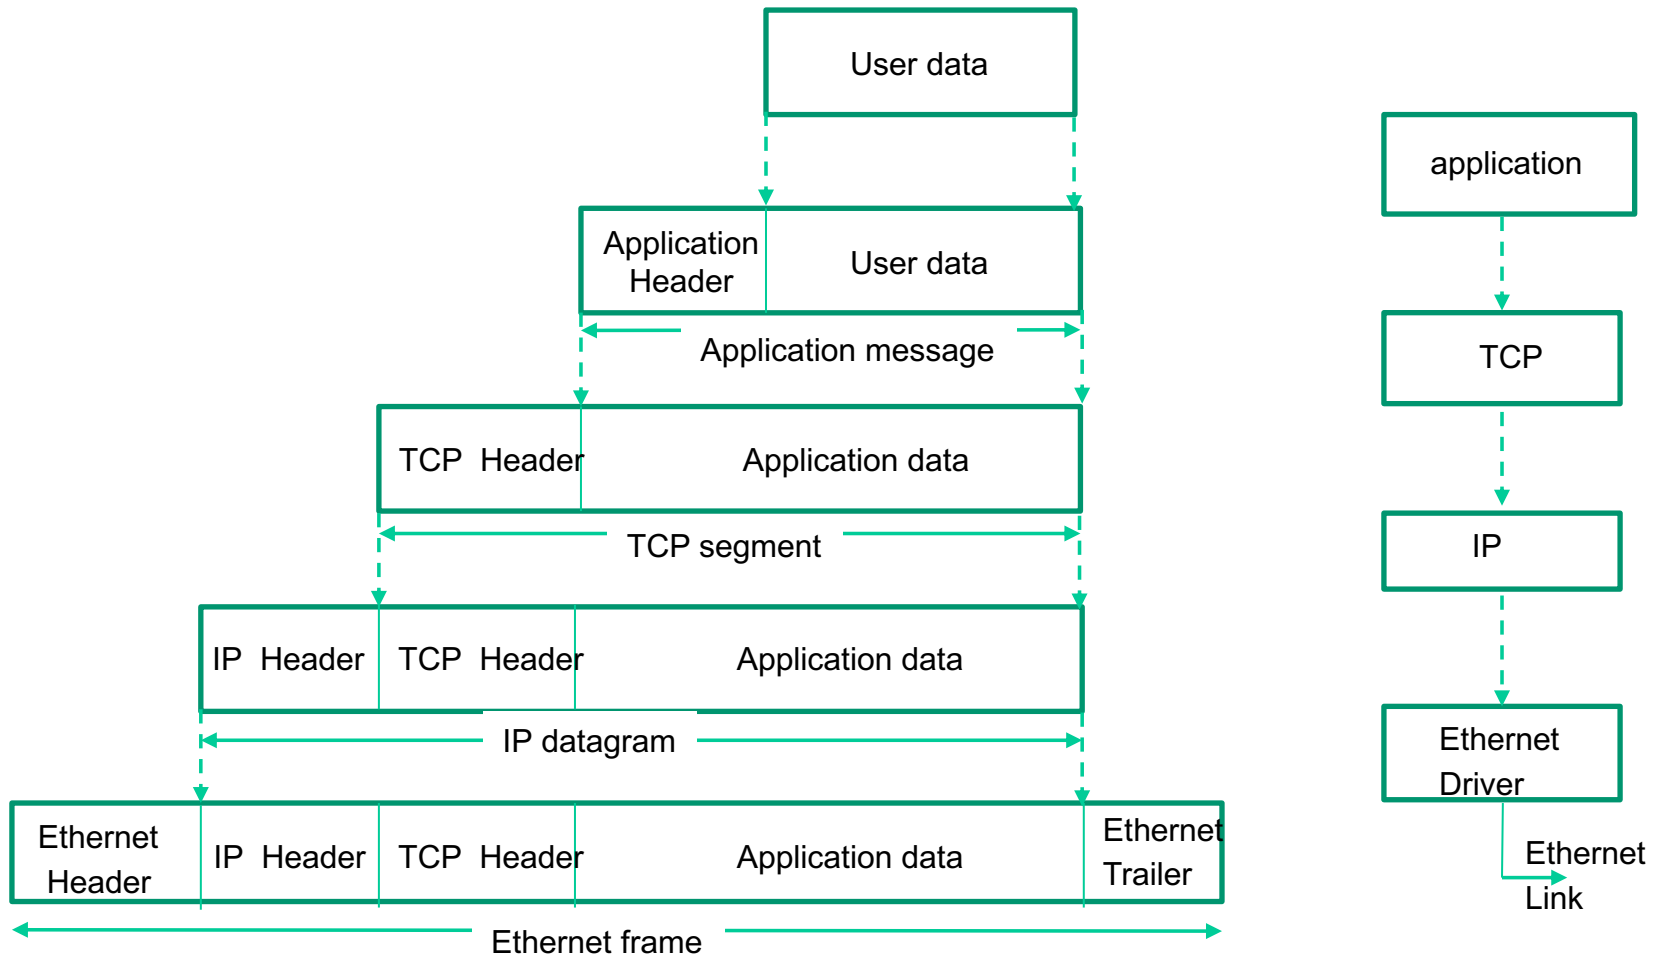
# RECAP: What is the Internet Architecture?

## Protocol Stack



# RECAP: What is the Internet Architecture?

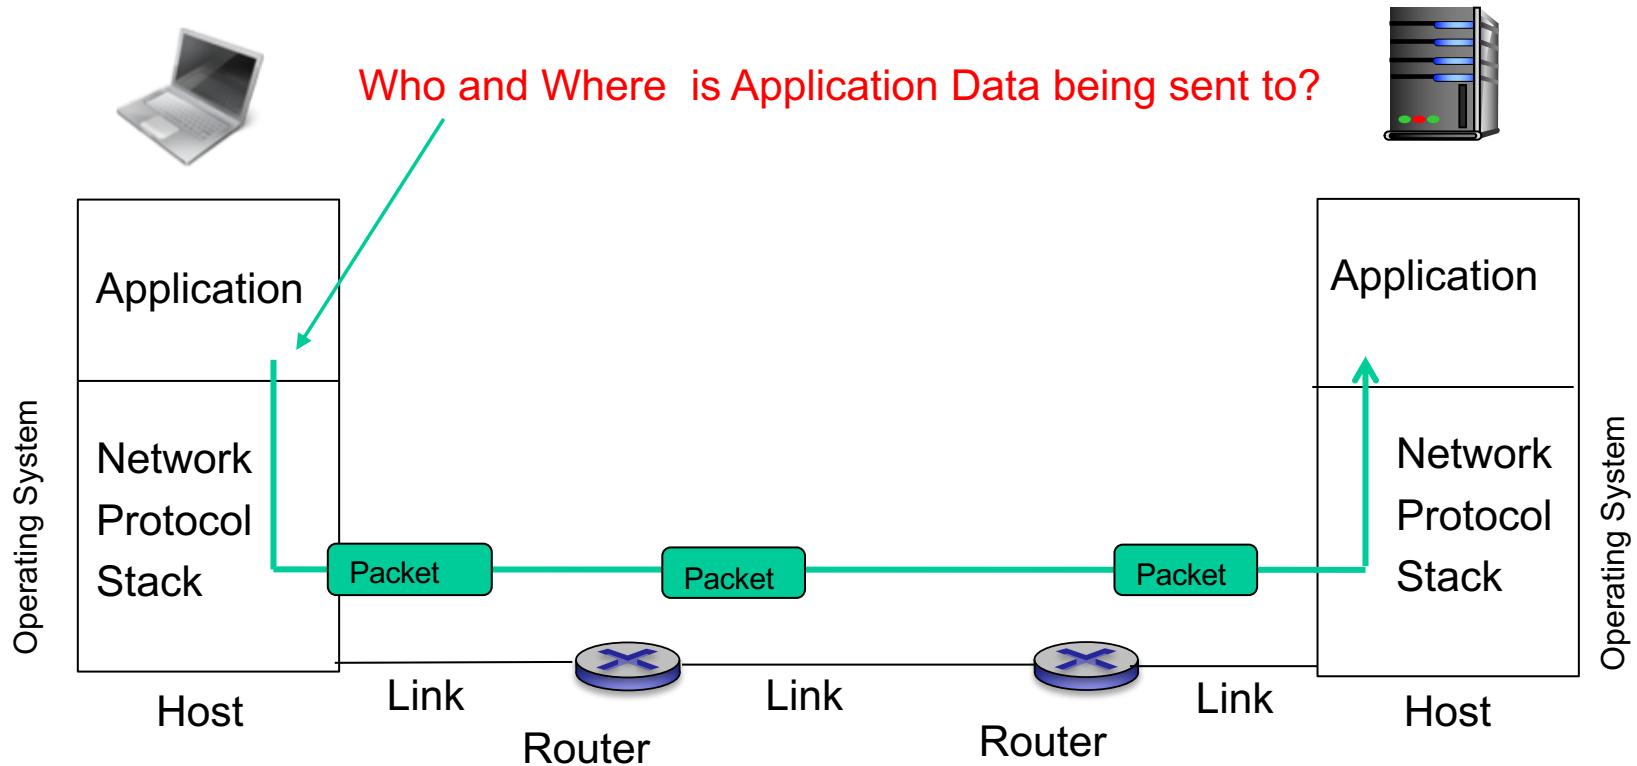
## Protocol Stack: Encapsulation



Reference: WR Stevens TCP Illustrated Volume 1, Addison-Wesley

# RECAP: What is the Internet Architecture?

## Packet Switching



Note: Uni-directional View





# RECAP: What is the Internet Architecture?

## Packet Switching



- Naming and Addressing the Destination App (Process)
  - IP address indicates destination host
  - Port indicates application on host
- Use Post Office Analogy

Letter



Postal Address identifies and locates house  
Name identifies person in house

IP Packet



**IP Address** identifies and locates host  
**Port** identifies application in host

# Host Addressing

- Hosts are addressed via **IP addresses** (32 bits)
  - Example: www.njit.edu = 128.235.251.25
  - IP addresses assigned by Regional Internet Registries (RIR)
  - To communicate with a destination, the application must determine the IP address of the host it is communicating with
    - May know a priori
    - May look up hostname->IP address mapping in DNS
    - May learn from incoming message (responses only)

# Process Addressing

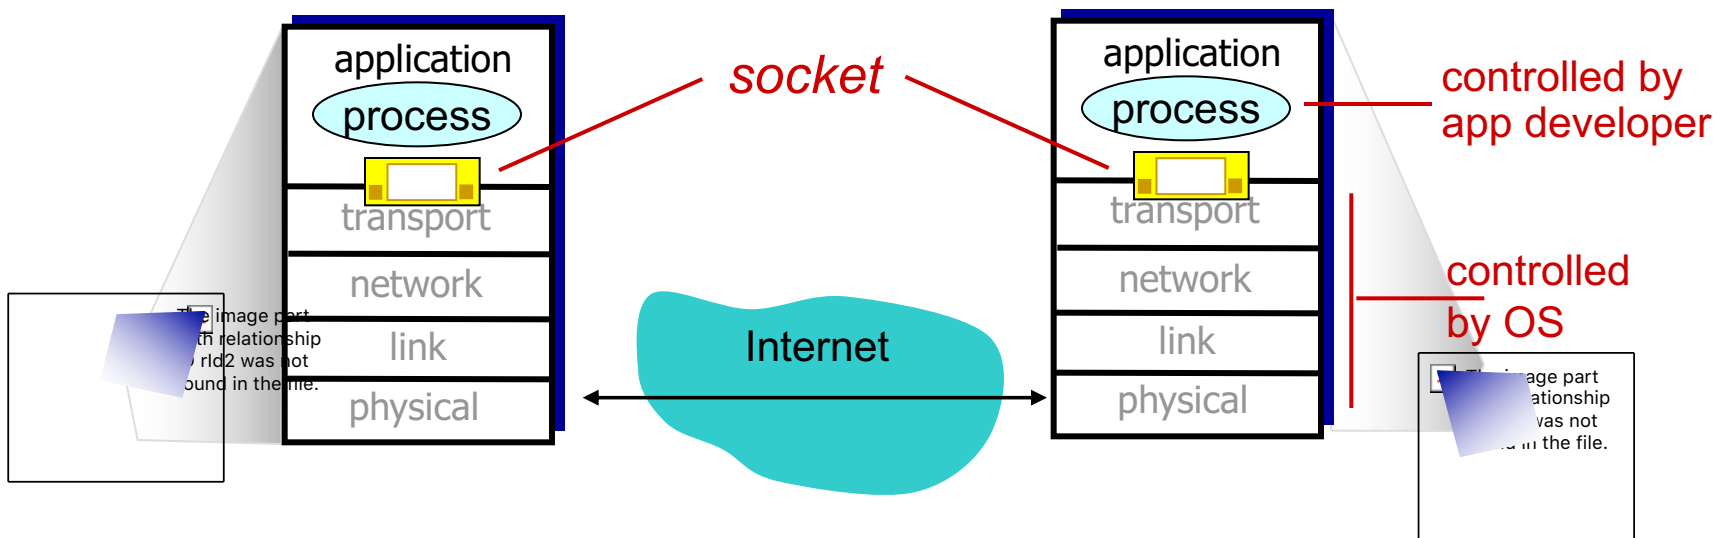
- **Ports** (16 bits)
  - Examples: 80 (HTTP), 53 (DNS)
  - Port numbers 0-1023 reserved for “well-known” services
  - Proprietary/Private apps: 1024 – 65535
  - Reserved Port Numbers assigned by IANA (Internet Assigned Numbers Authority) [www.iana.org](http://www.iana.org)
  - Operating system can allocate unique port number to application dynamically, or application developer can specify port number to use

# Socket API

The image part with relationship ID rld2 was not found in the file.

**goal:** learn how to build client/server applications that communicate using sockets

**socket:** door between application process and end-end-transport protocol




# Socket API

- Application developers need to send messages across the network
- The Socket API provides this:
  - Access to network controlled in OS
  - Other functions provided in system library

# Socket API

- Design Goals
  - Communication between processes should not depend on whether they are on same or different machine

# Socket API

 The image part with relationship ID rld2 was not found in the file.

*Two socket types for two transport services:*

- **UDP:** unreliable datagram
- **TCP:** reliable, byte stream-oriented

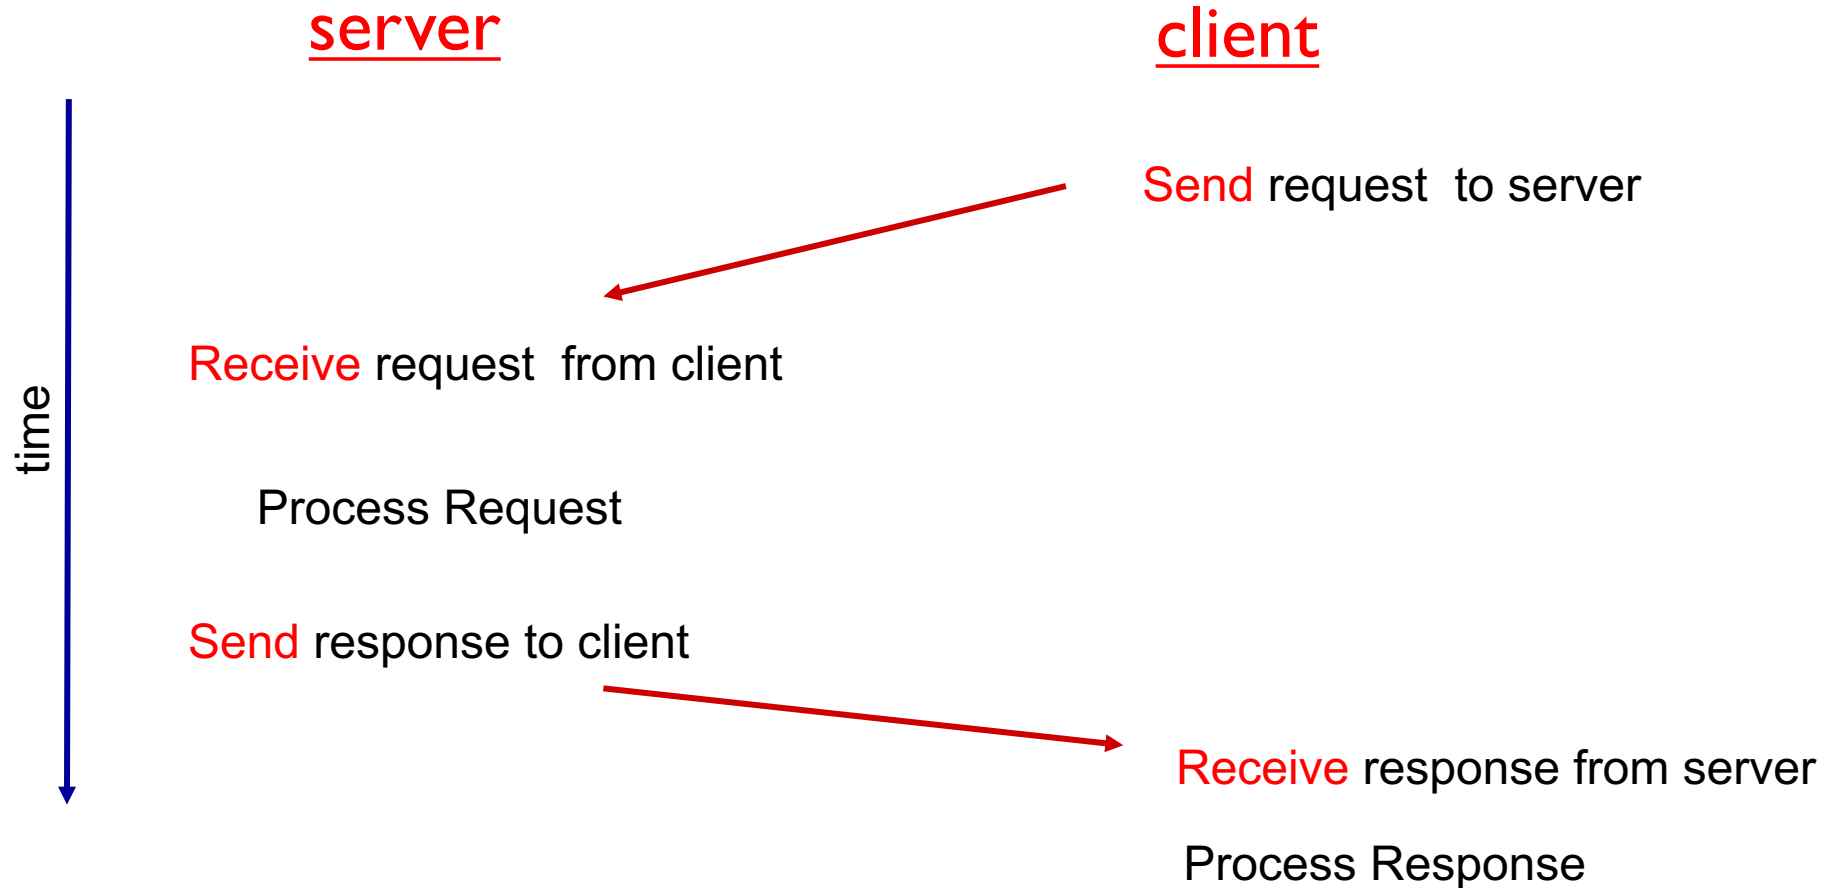
*Application Example:*

1. client reads arguments from command line and sends data to server
2. server receives the data and prints
3. server echoes received data back to client
4. client receives data and prints

# Client/server socket interaction:



The image part with relationship ID rld2 was not found in the file.





# Socket programming *with UDP*



The image part with relationship ID rld2 was not found in the file.

## UDP: no “connection” between client & server

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

## UDP: transmitted data may be lost or received out-of-order

### Application viewpoint:

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

# Client/server socket interaction (UDP):



The image part with relationship ID rld2 was not found in the file.

server

client

Create socket

Create socket

Bind socket  
(address, port)

Bind socket  
(address, port)

Receive message

Send message


Send message

Receive message

Close socket

Close socket

time



# Step 1

Create a socket

```
int s = socket (int family, int type, int protocol)
```

AF\_INET  
AF\_INET6

SOCK\_DGRAM (UDP)  
SOCK\_STREAM (TCP)  
SOCK\_RAW (IP)

Default = 0

Python:

```
clientSocket = socket(AF_INET,  
                      SOCK_DGRAM)
```

# Step 2


## Bind a socket

```
int error = bind (s, addr, addrlen)
```

socket # from  
*socket* call

Transport address  
(IP address, port)

Length of address



```
Python: serverSocket.bind("", serverPort)
```

# Step 3

- For TCP only, see later in slide deck
- Includes TCP connection establishment

# Step 4 (exchange data)

Python:

```
clientSocket.sendto(message.encode(),  
                    (serverName, serverPort))  
receivedMessage, serverAddress =  
    clientSocket.recvfrom(2048)
```

Connectionless calls (UDP):

```
int sendto(s, void *msg, int len, u_int flags,  
           struct sockaddr *toaddr, addrlen)  
int recvfrom (s, void *buf, int len, u_int flags,  
              struct sockaddr *fromaddr, addrlen)
```

# Step 5

Close socket

**close(s)**



Listening or connected socket #

# Client/server socket interaction: UDP



The image part with relationship ID rld2 was not found in the file.

## server (running on *serverIP*)

create socket, port= x:  
**serverSocket =**  
**socket(AF\_INET,SOCK\_DGRAM)**

↓  
read datagram from  
**serverSocket**

↓  
write reply to  
**serverSocket**  
specifying  
client address,  
port number

## client

create socket:  
**clientSocket =**  
**socket(AF\_INET,SOCK\_DGRAM)**

↓  
Create datagram with server IP and  
port=x; send datagram via  
**clientSocket**

↓  
read datagram from  
**clientSocket**

↓  
close  
**clientSocket**



# Example app: UDP client

## *Python UDPClient*

include Python's socket library

```
import sys, time
from socket import *
```

```
# Get the server hostname, port and data length as command line arguments
argv = sys.argv
host = argv[1]
port = argv[2]
count = argv[3]
```

```
# Command line argument is a string, change the port and count into integer
port = int(port)
count = int(count)
data = 'X' * count # Initialize data to be sent
```

create UDP socket for server

```
# Create UDP client socket. Note the use of SOCK_DGRAM
clientsocket = socket(AF_INET, SOCK_DGRAM)
```

Attach server ip, port to message; send into socket

```
# Sending data to server
print("Sending data to " + host + ", " + str(port) + ": " + data)
clientsocket.sendto(data.encode(),(host, port))
```

read reply characters from socket into string

```
# Receive the server response
dataEcho, address = clientsocket.recvfrom(count)
```

print out received string and close socket

```
# Display the server response as an output
print("Receive data from " + address[0] + ", " + str(address[1]) + ": " + dataEcho.decode())

#Close the client socket
clientsocket.close()
```

# Example app: UDP server

## *Python UDPServer*

```
import sys
from socket import *

serverIP = "          # any local IP address
serverPort = 12000
dataLen = 1000000

# Create a UDP socket. Notice the use of SOCK_DGRAM for UDP packets
serverSocket = socket(AF_INET, SOCK_DGRAM)
# Assign IP address and port number to socket
serverSocket.bind((serverIP, serverPort))

print('The server is ready to receive on port: ' + str(serverPort))

# loop forever listening for incoming datagram messages
while True:

    # Receive and print the client data from "data" socket
    data, address = serverSocket.recvfrom(dataLen)
    print("Receive data from client " + address[0] + ", " + str(address[1]) + ": " + data.decode())

    # Echo back to client
    print("Sending data to client " + address[0] + ", " + str(address[1]) + ": " + data.decode())
    serverSocket.sendto(data, address)
```

create UDP socket →

bind socket to local port  
number 12000 →

loop forever →

Read from UDP socket into  
message, getting client's  
address (client IP and port) →

send string back to this  
client →

# Socket programming *with TCP*



The image part with relationship ID rld2 was not found in the file.

## client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

## client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket:* client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients (more in Chap 3)

## application viewpoint:

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

# Client/server socket interaction (TCP):



The image part with relationship ID rld2 was not found in the file.

server

client

Create socket

(listening socket )

Create socket

Bind socket  
(address, port)

Bind socket  
(address, port)

Listen/ accept

(create new connection socket)

Connect to server

Read/write data

Read/write data

Close sockets

Close socket

time



# Step 1

Create a socket

```
int s = socket (int family, int type, int protocol)
```

AF\_INET  
AF\_INET6

SOCK\_DGRAM (UDP)  
SOCK\_STREAM (TCP)  
SOCK\_RAW (IP)

Default = 0

Python:

```
clientSocket = socket(AF_INET,  
                      SOCK_STREAM)
```

# Step 2


## Bind a socket

```
int error = bind (s, addr, addrlen)
```

socket # from  
*socket* call

Transport address  
(IP address, port)

Length of address



```
Python: serverSocket.bind("", serverPort)
```

# Step 3a (TCP server)

Listen on a socket

```
int error = listen (s, backlog)
```

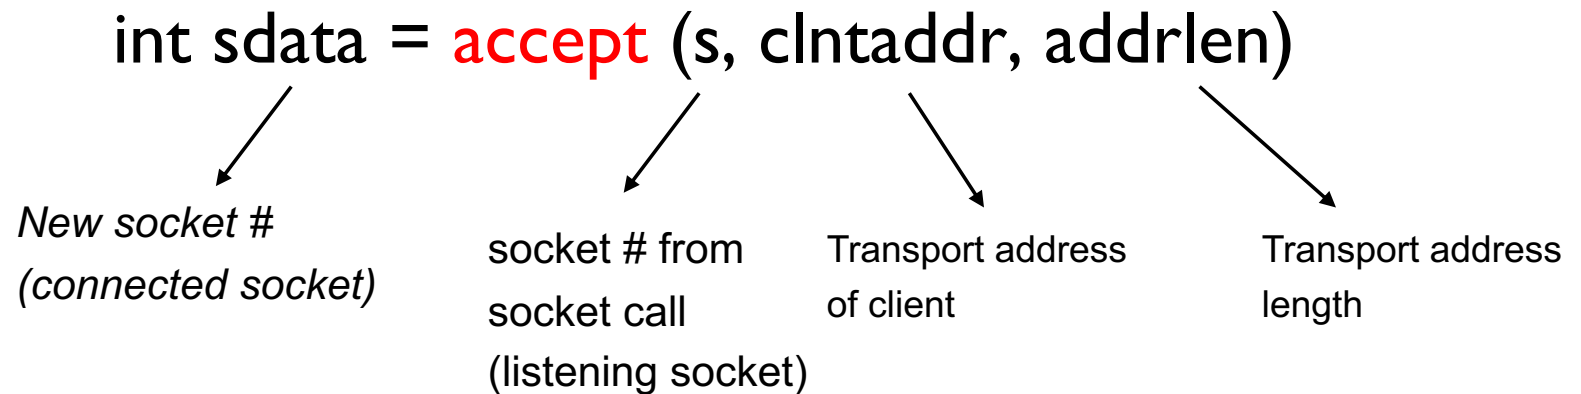
socket # from  
*socket* call

Queue Length  
(number of connections  
allowed between *accept*  
system call)

Socket now configured to listen to new connections. Data flows on another socket.

# Step 3b (TCP server)

Wait for new connection request from client



New socket created for data flow.



# Step 3 (TCP client)

Connect to server

`int error = connect (s, svraddr, addrlen)`

The diagram illustrates the parameters of the `connect` function call. Arrows point from the following text labels to their corresponding arguments in the code:

- New socket #* points to `s`.
- socket # from socket call points to `svraddr`.
- Transport address of server points to `svraddr`.
- Transport address length points to `addrlen`.

Client connection request to server

# Step 4 (exchange data)

Connection-oriented calls (TCP):

**read**, **write** system calls

int **send**(s, void \*msg, int len, u\_int flags)

int **recv**(s, void \*buf, int len, u\_int flags)

Connectionless calls (UDP):

int **sendto**(s, void \*msg, int len, u\_int flags,  
struct sockaddr \*toaddr, addrlen)

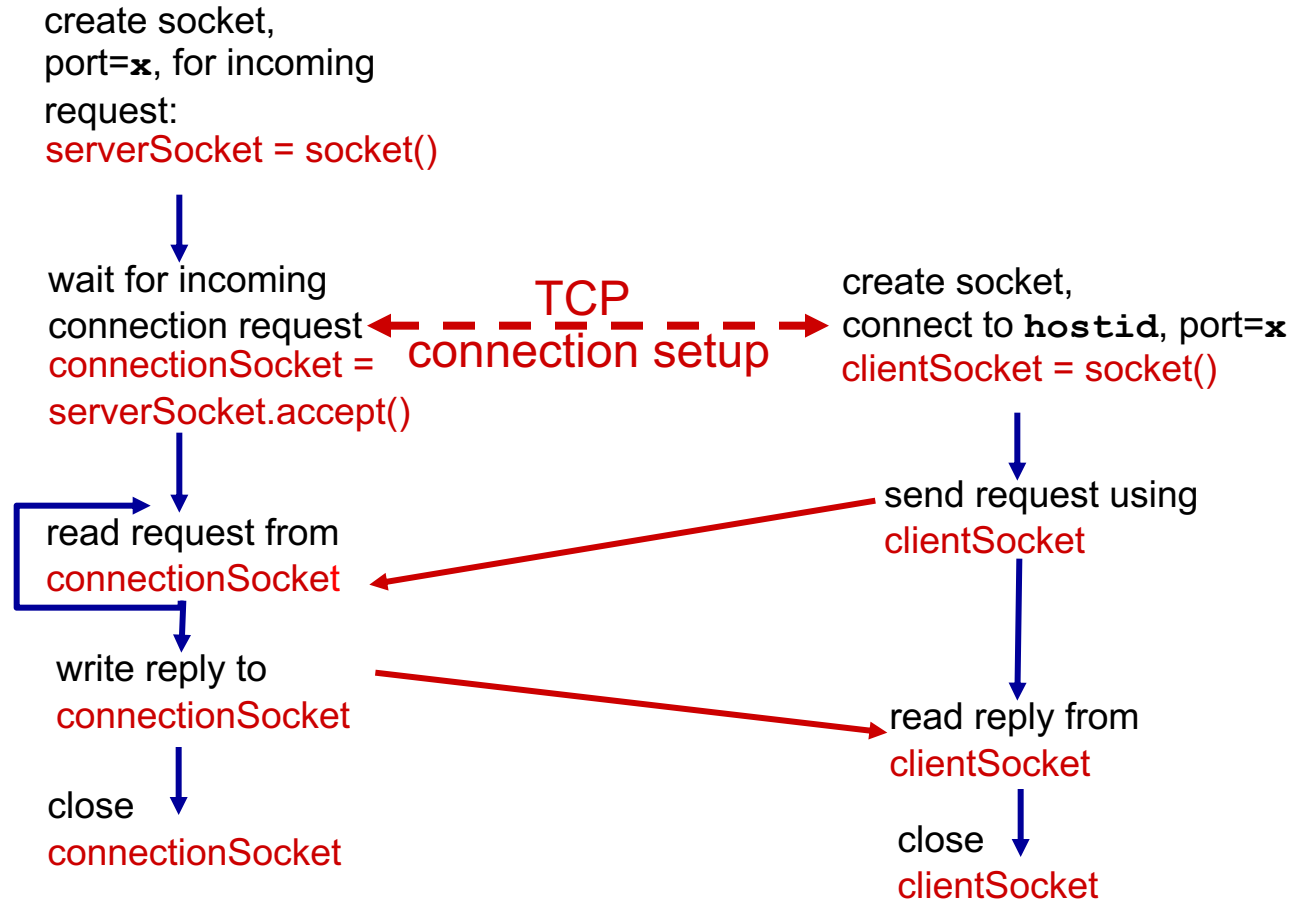
int **recvfrom** (s, void \*buf, int len, u\_int flags,  
struct sockaddr \*fromaddr, addrlen)

# Client/server socket interaction: TCP

The image part with relationship ID rld2 was not found in the file.

## server (running on `hostid`)

## client



# Example app: TCP client

create TCP socket and  
connect to server ip, port



No need to attach server ip,  
port



## *Python TCPClient*

```
import sys
from socket import *

# Get the server hostname, port and data length as command line arguments
argv = sys.argv
host = argv[1]
port = argv[2]
count = argv[3]

# Command line argument is a string, change the port and data length into integer
port = int(port)
count = int(count)

# Initialize and print data to be sent
data = 'X' * count

# Create TCP client socket. Note the use of SOCK_STREAM for TCP packet
clientSocket = socket(AF_INET, SOCK_STREAM)

# Create TCP connection to server
print("Connecting to " + host + ", " + str(port))
clientSocket.connect((host, port))

# Send data through TCP connection
print("Sending data to server: " + data)
clientSocket.send(data.encode())

# Receive the server response
dataEcho = clientSocket.recv(count)
# Display the server response as an output
print("Receive data from server: " + dataEcho.decode())

# Close the client socket
clientSocket.close()
```

# Example app: TCP server

## *Python TCPServer*

create TCP welcoming  
socket

server begins listening for  
incoming TCP requests

loop forever

server waits on accept()  
for incoming requests, new  
socket created on return

read bytes from socket (but  
not address as in UDP)

close connection to this  
client (but *not* welcoming  
socket)

```
from socket import *

serverIP = "          # any local IP address
serverPort = 12000
dataLen = 1000000

# Create a TCP "welcoming" socket. Notice the use of SOCK_STREAM for TCP packets
serverSocket = socket(AF_INET, SOCK_STREAM)
# Assign IP address and port number to socket
serverSocket.bind((serverIP, serverPort))
# Listen for incoming connection requests
serverSocket.listen(1)

print('The server is ready to receive on port: ' + str(serverPort))

# loop forever listening for incoming connection requests on "welcoming" socket
while True:
    # Accept incoming connection requests, and allocate a new socket for data communication
    connectionSocket, address = serverSocket.accept()
    print("Socket created for client " + address[0] + ", " + str(address[1]))

    # Receive and print the client data in bytes from "data" socket
    data = connectionSocket.recv(dataLen).decode()
    print("Data from client: " + data)

    # Echo back to client
    connectionSocket.send(data.encode())
    connectionSocket.close()
```

# POSIX system calls

## Summary

System Call	Description	
socket	Create a socket	
bind	Associate address, port with socket (sockaddr)	
listen	Set socket to listen to incoming connections	} only
accept	Accepts next client connection	
connect	Connect to socket on server	
Read/write, sendto, recvfrom Sendmsg, recvmsg	Exchange data	
Close/shutdown	Close connection	

server {

client {

Connection-oriented

# Byte Ordering

- An issue when sending multi-byte values
- Not all computers store bytes in multi-byte values in the same order
  - Little endian: low-order byte in address  $A$ , high-order byte in address  $(A+I)$
  - Big endian: high-order byte in address  $A$ , low-order byte in address  $(A+I)$
- Example: 4-byte integer  $0x00000001$  (value 1)

A	A+1	A+2	A+3
01	00	00	00

Little endian

A	A+1	A+2	A+3
00	00	00	01

Big endian

# Byte ordering: C

- Network protocols use big-endian ordering (\*)
- Conversion utilities available for programmer convenience

`uint_32 = htonl (uint32_t hostint32)`

`uint_16 = htons (uint16_t hostint16)`

`uint_32 = ntohl (uint32_t netint32)`

`uint_16 = ntohs (uint16_t netint16)`

*Reference: <https://linux.die.net/man/3/byteorder>*

*(\*) Includes when assigning IP address and port to `sockaddr_in` data structure*



# Byte ordering: Python


## ■ Python

- `struct.pack(format-string, v1, v2, ...)`
  - Return bytes object with values encoded according to format-string
  - Example: Encode unsigned integer `i` in network byte order in `msg`
    - `msg = struct.pack("!I", i)`
- `struct.unpack(format-string, buffer)`
  - Decode and return values packed in buffer according to format-string
  - Example:
    - `i = struct.unpack("!I", msg)`

*Python: See <https://docs.python.org/3/library/struct.html>*

*Java: Byte buffers use big-endian by default*

# Internet transport protocols services

 The image part with relationship ID rld3 was not found in the file.

## TCP service:

- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: OS maintains state

## UDP service:


- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

# When to use UDP instead of TCP

- Multicast or broadcast
- Simple request-response type applications
  - Small amount of data
  - Infrequent communications
  - Timeliness
  - But must implement reliability in application
  - For bulk transfers, use TCP (do not want to reinvent the wheel at application layer)

# What transport service does an app need?

 The image part with relationship ID rld3 was not found in the file.

## data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

## security

- encryption, data integrity, ...

# Transport service requirements: common apps



The image part with relationship ID rld3 was not found in the file.

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100' s msec yes and no

# Internet apps: application, transport protocols



The image part with relationship ID rld3 was not found in the file.

<b>application</b>	<b>application layer protocol</b>	<b>underlying transport protocol</b>
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP
Name Resolution	DNS	TCP or UDP
Host Configuration	DHCP	UDP

# References

## Network Programming and Socket API

### ■ Python

- <https://docs.python.org/3/library/socket.html>
- <https://docs.python.org/3/library/struct.html>

### ■ C

- <https://linux.die.net/man/7/socket>

### ■ Java

- <https://docs.oracle.com/javase/tutorial/networking/datagrams/index.html>
- <https://docs.oracle.com/javase/tutorial/networking/sockets/index.html>